

STRATEGY DESIGNED FOR HUMANS.

bp ST&S AI PoC — Capability & Resource Proposal

JUNE 2026



What we'll cover.

01 | The People in This Conversation

The team you'll be working with

02 | A New Kind of Partner

Built on 30 years of real delivery

03 | Your Ask

What we understood — and what we propose

04 | Discovery Agenda

The questions that shape the right engagement

05 | The Path Forward

Five steps from this conversation to day one

06 | This Is Harder Than It Looks

Why most AI deployments fail in production

07 | Our AI Track Record

Named engagements. Real outcomes. Delivered.

08 | Engineering for Determinism

Eight controls for trustworthy GenAI

09 | The Criticality Spectrum

How much rigour does your use case need?

10 | Risk Line of Sight

Traceability from risk to decision

11 | What We're Proposing

Closing summary

12 | Appendix

Technical reference material

The people in this conversation.



STEVEN WHALEY

Founder & CEO, Soda Partner

- Data and digital transformation
- Financial crime and regulatory change
- Complex program delivery
- Senior executive advisory

25+ years advising and delivering complex global transformation across Tier 1 and Tier 2 financial institutions.



JOHN ARGENT

Founder & CEO, Six Group

- Culture and human transformation
- Brand and organizational change
- Strategy activation
- Senior leadership engagement

30 years helping complex organisations make strategy stick.

Six Group — bp's contracted delivery partner



SCOTT ANDERSON

Co-Founder & Senior Partner, SODA Partner

- Enterprise Data & AI Strategy
- AI Use Case Identification & PoC Leadership
- Large-Scale Programme Delivery (\$100M+)
- LLM & Generative AI Applications

30+ years building and scaling data platforms and leading AI adoption across global financial institutions.



DRAZENKO IVEZIC

Co-Founder & Senior Partner, SODA Partner

- Financial Crime Data Strategy & Governance
- AI & Data Integrity Frameworks (TM, AML, Sanctions, CDD)
- Large-Scale Programme Delivery & Global Team Leadership
- Payments Infrastructure & Enterprise Data Platform Build

25+ years designing trusted data foundations and embedding practical AI and financial crime controls across Tier 1 global institutions.


Senior practitioners. Not account managers.
The people named here are the people who will work on your engagement.

A new kind of partner. Built on 30 years of real delivery.



THE SODA BELIEF:

- ✓ Human-led and outcome-driven
- ✓ Built with the people who will live with it
- ✓ Measured by capability gained, not hours billed



THE FOUNDATION (SIX GROUP):

- Backed by Six Group — a strategy activation partner since 1994
- **30+ years. 40+ clients.**
- Sectors: Financial Services, FMCG, Energy




S&P Global



SCOTTISH
WIDOWS

babcock™



New thinking. Proven heritage. One partnership model.

YOUR ASK

We heard you.

bp's Digital Innovation team needs production-grade AI capability, fast. Here is what we understood — and what we propose.

What we need to understand.

Use case first. Profiles second. Always.



USE CASE & TECHNICAL



Initial brief call with Angela Fine and Dan Golding, 28 May 2026.

Requirements document requested and pending receipt.



What is the specific trading use case?

Price signal modelling / trade documentation / risk classification / execution support? Understand Phase 1 PoC scope and what has been delivered. Confirm what Phase 2 is building on. Identify any broader interdependencies with other PoCs in the bp innovation backlog.



Internal or external data sources — or both?

Shapes the RAG architecture and data governance requirements from day one.



Cloud environment?

AWS / Azure / GCP — affects tooling and containerisation approach.



Current data foundations — quality, governance, lineage?

AI delivers value only when the data beneath it is sound.



What does production-grade look like for this PoC?

What is the success metric — and who signs it off?

**We ask these questions because we qualify fit
before we submit.**

From this conversation to **day one.**

Five steps. Clear ownership. No ambiguity.



THE SODA POSITIONING PLAY:

We are not a recruiter throwing CVs over a wall.

- We qualify fit before we submit — use case understanding first, profiles second
- We bring the advisory layer — determinism framework, governance thinking, AI system architecture
- We are not just the talent finder



95% of GenAI pilots fail to reach production. Our job is to be in the 5% that deliver on value.

OUR THINKING

This is harder than it looks.

Most AI deployments fail in production. Here is what it takes to be in the minority that don't — and why we know.

AI delivery. Proven.

Named engagements. Real outcomes. Delivered.



LLM-ENABLED DATA STRUCTURING

Leading Global Stock Exchange

Enterprise LLM deployed to extract and structure data from unstructured PDF documents. Integrated with structured database. Human-in-the-loop QA pipeline built. Reached production.

PRODUCTION DEPLOYED



DIGITAL TWIN & PREDICTIVE FAILURE MODELLING

Global Engineering Group

65 models built to digitally twin an industrial gas turbine — a critical asset in power generation. Predictive maintenance approach directly supports sustainability goals — reducing unplanned downtime and enabling more efficient, lower-emission power generation operations. Validated against blind field data — predicted actual turbine failure with 95% certainty. First known application of this approach in the industry.

95% ACCURACY



AI USE CASE DISCOVERY & PoC LEADERSHIP

Tier 1 Global Asset Manager

Led identification and business case development for AI use cases across customer 360, personalisation, segmentation, and cross-sell modelling. Outcome-led use case selection. Board-level investment secured.

BOARD INVESTMENT SECURED



ML FOR DATA QUALITY AUTOMATION

Global Custodian Bank

Designed ML approach to identify known data quality failure patterns and apply auto-correction at scale. Conceptual framework for self-healing data quality capability.

INNOVATION PROGRAMME



AI GOVERNANCE & TRUSTED AI AT SCALE

Major European Global Bank

Led design and challenge of AI-enabled Transaction Monitoring and Perpetual KYC capabilities supporting financial crime detection across millions of customers. Established traceability from risk typologies through customer behaviour, data features and AI outcomes, ensuring explainability, auditability and regulatory confidence.

REGULATED PRODUCTION ENVIRONMENT

150+ years of data and AI delivery across Tier 1 institutions. The model is not the difficult bit. Trusting the model is.

Engineering for Determinism — Eight Controls for Trustworthy GenAI in Production

Non-determinism is an engineering problem, not a fundamental barrier.
Build for it from day one.



KEY PRINCIPLE — The LLM is one component. The engineering around it is what makes it trustworthy.

1 TEMPERATURE = ZERO



Always select the highest-probability output. Lock this in every production system. No exceptions.

2 STRUCTURED OUTPUT SCHEMAS



Force responses into a fixed JSON format. Reject anything that does not conform. The model cannot wander.

3 RAG GROUNDING



Feed the model your verified data at inference time. It interprets facts — it does not generate them.

4 OUTPUT VALIDATION LAYERS



The model proposes. A deterministic rule layer disposes. Nothing reaches the user unvalidated.

5 DATA INTEGRITY ASSURANCE



Before model execution: completeness checks. Quality validation. Lineage verification. Drift monitoring. Perfect prompts on corrupted data still produce unreliable outcomes.

6 PROMPT ENGINEERING DISCIPLINE



Constrained, precise prompts reduce variance at source. Role. Task. Example. Constraints.

7 HUMAN-IN-THE-LOOP



AI narrows the decision space. A qualified human confirms. The audit trail belongs to a person.

8 SEMANTIC CACHING



Serve validated outputs for known query patterns. The model only runs for genuinely novel inputs.

JP Morgan. Goldman Sachs. Bloomberg. The firms getting this right are not using better models but are designing with this framework engineered in.

How much **determinism** does your use case need?

Applying GenAI engineering controls proportionately to use case risk

LOW CRITICALITY

CRITICAL

	LOW Document summarisation, content drafting	LOW-MEDIUM Customer chatbot, knowledge search	MEDIUM Research synthesis, contract review	MEDIUM-HIGH Compliance monitoring, risk reporting, regulatory horizon scanning	CRITICAL Trading signal support. Position risk classification. Real-time market pricing. Regulatory filings.
1. Temperature = Zero	●	●	●	●	●
2. Structured Output Schemas	○	●	●	●	●
3. RAG Grounding	○	○	●	●	●
4. Output Validation Layers	○	○	●	●	●
5. Data Integrity Assurance	○	○	●	●	●
6. Prompt Engineering Discipline	○	○	○	●	●
7. Human-in-the-Loop	○	○	○	●	●
8. Semantic Caching	○	○	○	○	●



bp ST&S Trading PoC

Commodities trading.
Real-time. Commercially
consequential. Regulated.

Full engineering stack
required from the first
line of code.




The cost of building these controls in from the start is low.
The cost of discovering you need them after a production failure is not.

Why should we trust the answer?

In regulated environments, traceability is not optional — it is the architecture.

The principle



In a trading context



In regulated environments, every AI-assisted decision must be traceable back to a risk typology. The chain of accountability is not optional — it is the architecture.

**If you cannot trace the decision back to the risk,
you cannot defend it to a regulator.**

What we're **proposing.**



WE UNDERSTOOD YOUR BRIEF

bp's innovation function needs specialist AI/ML engineers with production-grade delivery experience, fast. We have the people, the network, and the qualification process to put the right names in front of you — not a shortlist of CVs.



The brief is heard.



WE BRING MORE THAN RESOURCE

SODA brings an advisory layer alongside every engagement. Engineering for determinism. Risk line of sight. Data integrity. The framework that separates a successful PoC from a production system that actually works.



The expertise is here.



HERE IS THE PATH

Requirements doc reviewed. Profiles submitted within 24 hours. Candidates prepared. Mobilisation targeting week commencing 8th June. Five steps. Clear ownership. No ambiguity.



The path is clear.

STRATEGY DESIGNED FOR **HUMANS.**

SODA Partner × SIX Group

Appendices.

01

The Anti-Hallucination Discipline

How prompt-level verification maps to production determinism

02

From Sandbox to Production

The GenAI deployment pipeline from experimentation to live

The Anti-Hallucination Discipline – How It Maps to Production Determinism

Key components of a structured AI verification framework — and which of the eight controls each one addresses

KEY INSIGHT — The anti-hallucination framework is prompt engineering discipline (Control 6) applied specifically to factual accuracy. It operationalises Controls 3, 4, and 7 at the point of generation.

ANTI-HALLUCINATION COMPONENTS			MAPS TO WHICH CONTROL		THE 8 PRODUCTION CONTROLS	
1	Search Before Claiming	Verify every technical fact against current official sources before stating it	→	Controls 3 + 4	1	Temperature = Zero
2	Source Hierarchy	Official docs first. Community last. Recency matters.	→	Control 3	2	Structured Output Schemas
3	No Training Data Alone	Never answer purely from training data. Always verify.	→	Control 3	3	RAG Grounding
4	Confidence Scoring	Surface uncertainty when below threshold. Flag it — don't fill it with a guess.	→	Controls 4 + 7	4	Output Validation Layers
5	Anti-Confabulation Rules	Never invent parameters, feature names, or capability claims. If unverifiable, say so.	→	Controls 2 + 4	5	Data Integrity Assurance
6	Not Documented Protocol	State clearly when a feature cannot be verified. Label speculation explicitly.	→	Control 7	6	Prompt Engineering Discipline
7	Recency Gate	Flag sources older than 3 months for fast-moving tools. Freshness is a quality signal.	→	Control 3	7	Human-in-the-Loop
8	Source Integrity Validation	Trusted sources can still contain incomplete, delayed or incorrect data. Source validation should assess quality, completeness and consistency — not just provenance.	→	Controls 3 + 4 + 5	8	Semantic Caching

Controls 1 (Temperature) and 8 (Semantic Caching) operate at the infrastructure layer. Control 5 (Data Integrity Assurance) operates at the data layer, before model execution. Anti-hallucination discipline applies primarily at the generation and validation layers (Controls 2–4, 6–7).



Prompt-level verification is the first line of defence. It does not replace downstream validation — it reduces the error rate before validation is even needed.

From Sandbox to Production — The GenAI Deployment Pipeline

How responsible AI systems are built, tested, governed, and continuously improved



VERSION CONTROL (GITHUB) —
All changes versioned, reviewed, and promoted deliberately.

Code PR

Prompt PR

Schema change

Model version pin

Production tag

1 SANDBOX

- Prompt exploration.
- Agent design.
- RAG pipeline on sample data.
- Temperature variable.
- No live data.
- No commercial consequence.

Gate:

Does the prototype produce plausible structured outputs against sample data?

2 PoC / PROTOTYPE

- Defined use case.
- Architecture decisions.
- Prompt engineering formalised.
- Schema defined.
- Small user group.

Gate:

Does the PoC demonstrate the use case is viable and data foundations sufficient?

3 DEVELOPMENT & TESTING

- Full architecture build.
- RAG on representative data.
- Validation layers built.
- HITL review process defined.
- Quality gate test suite.
- Temperature locked.

BRANCH:

Single LLM (e.g. Claude/Anthropic — pinned version) **OR** Multi-LLM (model router — agnostic layer, each model tested independently)

Gate:

Does the system pass all quality gates? Schema validation rate? HITL agreement rate?

4 STAGING

- Representative live-equivalent data.
- Full quality gates running.
- Security and compliance review.
- Model version pinned.
- Audit logging active.

Gate:

Does the system meet accuracy, latency, and compliance thresholds? Sign-off from risk/compliance/engineering.

5 PRODUCTION

- Real data. Real users.
- Real consequences.
- All eight controls active.
- Every output logged — model version, prompt version, RAG source version, timestamp, HITL reviewer.
- Continuous monitoring for drift.
- Drift detection triggers staging review.
- Never auto-upgrade.

CONTINUOUS IMPROVEMENT LOOP



HITL disagreements → retraining signal



Drift detection → prompt review



Model provider update → controlled promotion



User feedback → backlog



A production GenAI system is not a model with a user interface. It is a layered architecture where the LLM is one component — surrounded by the engineering that makes it trustworthy.